

Database - P93

Redis

APPEND <key> <value> : append value to key
if <key> not exist, create and set as empty string

STRLEN <key> : get length of value stored in key
return error if <key> hold non-string value

although store string, recognize integer and provide simple operation

INCR <key> : increment integer value of key by one

INCRBY <key> <increment> : increment integer value of key by given amount

INCRBYFLOAT <key> <increment> : increment float value of key by given amount

DECR <key> : decrement integer value of key by one

DECRBY <key> <decrements> : decrement integer value of key by given amount

if <key> not exist, set to 0 before performing operation

return error if <key> contain wrong type value or string not represented as integer
limit to 64-bit signed integer

GETSET <key> <values> : set string value of key and return old value

return error when <key> exist but not hold string value

MULTI block atomic commands : similar concept as transaction support in other database
wrapping operations in single block will complete either successfully or not all
will never end up with partial operation

MULTI : make the start of transaction block

EXEC : execute all commands issued after MULTI

DISCARD : flush all previously queued commands in transaction

and restores connection state to normal

similar to ROLLBACK in SQL, clear transaction queue

unlike ROLLBACK, will not revert the database

simply will not run transaction at all

end effect identical, while underlying concept different

transaction rollback - operation cancellation

WATCH [<key>] : mark given key to be watched for conditional execution

UNWATCH : flush all previously watched key for transaction

no need to manually call if call EXEC / DISCARD

Database - Pg 4

Redis

collection datatype can contain huge number (up to 2^{32}) of values per key
command pretty to decode with the pattern

set command begin with S

hash command begin with H

sorted set command begin with Z

list command begin with L (for left) / R (for right)

depending on direction of operation

Hashes: like nested Redis object that can take any number of key-value pairs

avoid storing data with artificial key prefix

colon ":" for logical separation of key into segments

only matter of convention

HSET <key> <field> <value> : set string value of hash field

HMSET <key> [<field> <value>]⁺ :

set multiple hash fields to multiple values

HSETNX <key> <field> <value> :

set value of hash field only if field not exist

HGET <key> <field> : get value of hash field

HMAET <key> [<field>]⁺ : get values of all given hash fields

HGETALL <key> : get all fields and values in hash

HKEYS <key> : get all fields in hash

HVALS <key> : get all values in hash

HLEN <key> : get number of fields in hash

HSTRLEN <key> <field> : get length of value of hash field

(value): return nil when <field> not present or <key> not exist

(length): return 0 when <field> not present or <key> not exist

HEXISTS <key> <field> : determine if hash field exist

HDEL <key> [<field>]⁺ : delete given hash field

HINCRBY <key> <field> <increment> : increment integer value of hash field by given number

HINCRBYFLOAT <key> <field> <increments> :

increment float value of hash field by given amount

Database - P95

Redis

List : contain multiple ordered values

can act both as queue (FIFO) and as stack (LIFO)

more sophisticated actions

inserting somewhere in the middle / constraining size / moving between lists

Redis command return number of value pushed when inserting

list operation in Redis use zero-based index position

(a negative position : number of step from the end)

LPUSH <key> [<value>]+ : prepend one or more values to list

element inserted one after other to the head of list

RPush <key> [<value>]+ : append one or more values to list

element inserted one after other to the tail of list

if <key> not exist, created as empty list before performing push operation

Lpushx <key> <value> : insert value at the head of list

Rpushx <key> <value> : insert value at the tail of list

only if <key> already exist and hold list

Lset <key> <index> <value> : set value of element in list by index

return error for out of range indexes

Linsert <key> BEFORE | AFTER <pivot> <value>

insert element before/after another element in list

Llen <key> : get length of list

Lindex <key> <index> : get element from list by index

negative indices can be used to designate element starting at tail

-1 : last element, -2 : penultimate element, ...

return nil when index out of range

Lrange <key> <start> <stop> : get range of elements from list

out of range indexes will not produce error

Lpop <key> : remove and return first element of list

Rpop <key> : remove and return last element of list

Rpoplpush <source> <destination> :

return and remove the last element of <source>

push the element at the first element of <destination>

Database - Pg 96

standard

Pg 97 - command

Redis

List : LTRIM <key> <start> <stop> : trim list to specified range
out of range indexes will not produce error

useful when Redis to store logs

LREM <key> <count> <value> : remove elements from list

<count> = 0 : remove the first <count> occurrences of elements equal to <value>

<count> = 0 : remove all elements equal to <value>

<count> > 0 : remove elements equal to <value> from head to tail

<count> < 0 : removes elements equal to <value> from tail to head

return number of removed elements

blocking list : BLPOP [<key>] + <timeout>

BRPOP [<key>] + <timeout>

remove and get first/last element in list

or block until one available

blocking version of LPOP / RPOP

block connection when no element to pop from any of given lists
element popped from head/tail of first non-empty list

return multi-bulk (<key>, <value>)

nil multi-bulk when no element popped and <timeout> expired

BRPOPLPUSH <source> <destination> <timeout>

blocking variant of RPOPLPUSH

<timeout> = 0 : used to block indefinitely

behave exactly like LPOP / RPOP / RPOPLPUSH

return nil when used inside MULTI / EXEC block

Set : SADD <key> [<member>] + add one or more members to set

SCARD <key> : get cardinality of set

SMEMBERS <key> : get all members of set

SISMEMBER <key> <members> : determine if given value is member of set

SMOVE <source> <destination> <members> :

<members> removed from <source> and add to <destination> if exist

appear to be member of <source> or <destination>

for other clients in every given moment

Database - Pg 7

Redis

Redis

Set: `SUNION [<key>] +`: union multiple sets
`SINTER [<key>] +`: intersect multiple sets
`SDIFF [<key>] +`: subtract multiple sets
左结合的: `SDIFF A B C` 等价于 $(A \setminus B) \setminus C$
`SUNIONSTORE <destination> [<key>] +`
`SINTERSTORE <destination> [<key>] +`
`SDIFFSTORE <destination> [<key>] +`
union / intersect / subtract multiple sets and store resulting set
overwrite `<destination>` if exist

`SRANDMEMBER <key> [<count>]`
get one or more random members from set
`<count>` omitted: return one random element from set
`<count>` positive: return array of `<count>` distinct elements
`<count>` negative: return array of `<count>` elements
may return same element multiple times

`SPOP <key> [<count>]`
remove and return one or more random members from set
`SREM <key> [<members>] +`
remove specified members from set
treat `<key>` as empty set if not exist

unlike list, no blocking command for set

sorted set: ordered like list, but unique like set
field-value pair like hashes, but field: numeric score and denote

as random access priority queue

`ZADD <key> [NX|XX] [CH] [INCR] [<score> <member>] +`

add one or more members to sorted set + or update score if exist

`[<score> <member>] +` = score value: string representation of double precision floating point number

`NX|XX`: always/never add new member, not/only update member exist

`CH`: return number of new member added and member score updated

`INCR`: act like `ZINCRBY`, only one score-member can be specified

Database - Pg8

pp9 - sorted set



Redis

sorted set : `ZCARD <key>` : return cardinality of sorted set

~~signature~~ `ZSCORE <key> <member>` : get score associated with given member

~~return double~~ precision floating point number represented as string

`ZRANK <key> <member>` : determine index of given member

~~return index of member from scores ordered from low to high (ascending)~~

~~rank/index is 0-based~~

`ZREVRANK <key> <member>` :

~~get rank of member with scores ordered from high to low (descending)~~

~~return nil if <member> not exist / <key> not exist~~

`ZCOUNT <key> <min> <max>`:

~~count members with scores within given values~~

`ZINCRBY <key> <increment> <member>`

~~increment score of given member by given increment~~

~~if <member> not exist, added with <increment> as score~~

~~if <key> not exist, new sorted set created with specified <members>~~

~~Example 2 (Part 2 of 2)~~

`ZRANGE <key> <start> <stop> [WITHSCORES]`:

~~return range of members in sorted set by index~~

~~member considered to be ordered in ascending~~

`ZREVRANGE <key> <start> <stop> [WITHSCORES]`:

~~member considered to be ordered in descending~~

~~<start>, <stop> are inclusive ranges~~

~~the former result to now~~ `WITHSCORES` : return scores of members together with members

`ZRANGEBYSCORE <key> <min> <max> [WITHSCORES] [LIMIT <offset> <count>]`

`ZREVRANGEBYSCORE <key> <min> <max> [WITHSCORES] [LIMIT <offset> <count>]`

~~return range of members by score~~

`ZRANGEBYLEX <key> <min> <max> [LIMIT <offset> <count>]`

`ZREVRANGEBYLEX <key> <min> <max> [LIMIT <offset> <count>]`

~~adjustable in the range <min>, <max>: (/) specify range item exclusive/inclusive~~

~~use + to denote does not want to include + / - special meaning or positive/negative infinite string~~

~~if an item is used without plus sign then tinf / -inf no required to know highest/lowest score~~

`LIMIT` : only get range of matching member

~~similar to SELECT LIMIT offset, count in SQL~~

Database - Pg 9

8P9 - sorted set

Redis

sorted set: `ZLEXCOUNT <key> <min> <max>`

return count number of members between given lexicographical range
range can be in descending order when all elements inserted with same score

remove a range to zremrange <key> <start> <end>

`ZREM <key> [<member>]+`: remove one or more members from sorted set

`ZREMRANGE BY RANK <key> <start> <stop>`:

remove all members within given index

`ZREMRANGE BY SCORE <key> <min> <max>`:

remove all members with score between <min> and <max> (inclusive)

`ZREMRANGEBYLEX <key> <min> <max>`:

remove all members between given lexicographical range

when all elements inserted with same score

remove a range to zremrange <key> <start> <end>

remove a range to zpopmax <key> [<count>]

`ZPOPMIN <key> [<count>]`

remove and return up to count members with highest/lowest scores

<count> higher than cardinality will not produce error

one with highest/lowest will be the first when returning multiple elements

`BZPOPMAX <key>+ <timeout>`

`BZPOPMIN <key>+ <timeout>`

blocking variant of ZPOPMAX / ZPOPMIN primitive

block connection when no member to pop from any of given sorted set

`[<key>+<key>+TIME] [ZUNIONSTORE] <destination> <numkeys> [<key>]+`

`[WEIGHTS [<weight>]+] [AGGREGATE SUM|MIN|MAX]`

`[<key>+<key>+TIME] [ZINTERSTORE] <destination> <numkeys> [<key>]+`

`[WEIGHTS [<weight>]+] [AGGREGATE SUM|MIN|MAX]`

compute union/intersection of numkeys sorted set and store result in destination

possible to specify multiplication factor for each sorted set

score returned will be its score of every element multiplied by factor before AGGREGATE

`AGGREGATE`: possible to specify how result of union/intersection aggregated

`[<key>+<key>+TIME] [ZRETCOUNT] <numkeys>`

Database - P100

Redis

fast-access cache for data expensive to retrieve or compute
common use case for key-value system
ensuring content expire after designated time period
essential function to keep the set from growing unboundedly

EXPIRE <key> <seconds> : set key's time to live in seconds

PEXPIRE <key> <milliseconds> : set key's time to live in milliseconds
key automatically deleted after timeout expired

key with associated timeout is volatile in Redis terminology

EXPIREAT <key> <timestamp> : set expiration of key as UNIX timestamp

PEXPIREAT <key> <milliseconds-timestamp> : set expiration for key as UNIX timestamp specified in milliseconds

absolute UNIX timestamp since January 1, 1970

key deleted rather than expired

if EXPIRE/PEXPIRE with non-positive timeout

or EXPIREAT/PEXPIREAT with timestamp in the past

return 1 if timeout set, 0 if key not exist

EXISTS <key>^+ : return if key exist

also will return total number of given keys existing in item database

TTL <key> : get time to live for key in seconds

PTTL <key> : get time to live for key in milliseconds

return -1 if key exist but no associated expire

return -2 if key not exist

PERSIST <key> : remove expiration from key

turn key from volatile to persistent

return 0 if key not exist / no associated timeout

DEL <key>^+ : remove specified key

update expire time whenever retrieve value

common trick for keeping only recently used key

referred to most recently used (MRU) caching algorithm

unused key just expire as normal

Database - P101

Redis - command reference

Redis

not support expiring individual item in collection

EXPIRE keyword only work on top-level key

(ex) provide sufficient functionality in sorted set to accomplish

high level library (& Redisson for Java) provide functionality as well

repurpose score of item inserted into sorted set

to be expiration time (UNIX timestamp)

so ZADD jobs (now) + 60 job1

priority equal to timestamp of 60s from now

future request can quickly determine from the set

cron job set up to expire data easily from the set

so ZREM RANGE BYSCORE jobs 0 now)

namespace: referred to database, keyed by number

default database known as database 0

SELECT <index>: change selected database for current connection

(new, having specified 0-based numeric index)

new connection always use database 0

all databases anyway persisted together in same RDB/AOF file

different databases can have key having same name

MOVE <key> <db>: move key to another database

do nothing if <key> already exist in destination database

or not exist in source database

possible to use as locking primitive

useful for different applications running against single Redis server

need to trade data between each other

SWAPDB <index> <index>: swap two Redis database

immediately all clients connected to given database

will see data of other database

FLUSHDB [ASYNC]: delete all keys of currently selected DB

never fail: don't do this in production

RANDOMKEY: return random key from currently selected database

RENAME <key> <newkey>: rename key to new key

return error if <key> not exist

<newkey> overwritten if already exist

Database - P102

smithbala

3.9 - pipelined

Redis

telnet : Streaming commands through TCP directly interact without command-line interface terminating command with CRLF (carriage return/line feed, ~~回车/换行~~) output returned different to on the console

Redis streams OK status prefixed by "+"

number prefixed by "\$": number of characters in following string

number prefixed by ":" : number of values added successfully

"*" and number : number of complex values about to be returned

pipelining : netcat (nc) : BSD (Berkeley Software Distribution / Berkeley UNIX) command

Stream own string one at a time

must specifically end line with CRLF

pipelining / streaming multiple commands into single request

should always be preferred if make sense

\$0 echo -en "PING\r\nPING\r\nPING\r\n";

(ctrl+D) + cat = (sleep 1) | nc localhost 6379

pub-sub channel

publish-subscribe : several subscribers want to read announcements of single publisher

SUBSCRIBE [<channel>] +

listen for message published to given channels

cause CLI to output "Reading message ..."

block while subscriber listen for incoming message

not supposed to issue any other commands / PUBLISH

except for additional SUBSCRIBE/PUBLISH/UNSUBSCRIBE

PSUBSCRIBE [<pattern>] +

listen for message published to channel matching given patterns

glob-style pattern : ? : any single character

* : zero or more any characters

[..] : character in list

\ : to escape special character

UNSUBSCRIBE [<channel>] # +

UNPSUBSCRIBE [<pattern>] # +

unsubscribe client from given channels / patterns,

if none given, from all of ~~channels~~ channels / patterns

Database - P103

LPT - pending

Redis publish-subscribe : PUBLISH <channel> <message> : post message to given channel
return number of clients received message
all subscribers receive multibulk reply :
string "message"/ channel name / published message value
have to CTRL-C to break connection on redis-cli

server : INFO [<section>] : get information and statistics about server
<selection> : select specific section

general : general information about Redis server

clients : client connections section

memory : memory consumption related information

persistence : RDB and AOF related information

stats : general statistics

replication : master/replica replication information

CPU : CPU consumption statistics

commandstats : Redis command statistics

cluster : Redis cluster section

key space : database related statistics

all : return all sections

default : return only default set of section

durability : no persistence : Redis keep all value in main memory
good choice for basic cache server
durability increase latency

Redis built-in support for storing value to disk

rather than other fast-access cache

SAVE : perform synchronous save of dataset

producing point in time snapshot of all data in Redis instance

dataset snapshot in form of RDB file

block all other clients in production environment

BGSAVE : asynchronously save dataset to disk (in background)

LASTSAVE : get UNIX time stamp of last successful save to disk

not useful for check if BGSAVE command succeeded

this can return just not true

Database - P104

Redis

snapshotting : alter the rate of storage to disk

by adding / removing / altering one of save fields
each prefixed by "save" keyword

followed by time in second

and minimum number of keys must change before write occurs

save 900 : save 900 if any key changes saved in at least 900s

save 60 10000 : save 60 if 10000 keys change, save in 60 second

append-only file : Redis eventually durable by default

asynchronous write value to disk in interval defined

can also be forced to write by client-initiated command
acceptable for scenario : second-level cache / session server

insufficient for durable thing, like financial data

appendonly.aof : append-only file provided

keep record of all write command

like write-ahead log in HBase

"appendonly" must be enabled by setting to "yes" to redis.conf

(slow) always : more durable because every command saved

(default) everysec : save update/write command only once a second

no : let OS handle flushing

security : not natively built to be fully secure server

AUTH <password> : request for authentication in password-protected server

can be instructed to require password before allowing execute command

done using "requirepass" directive in configuration file

can be safely ignored because merely scheme for setting plaintext password

also provide command-level security through obscurity

\$O rename-command FLUSHALL (283d93ac95...)

相当于用 (283d93ac95...替換命令 FLUSHALL

configuration

can disable command entirely by setting to blank string in

\$O rename-command FLUSHALL ""

can set any number of commands to blank string

allow modicum of customization over command environment

Database - P105

Redis replication : Redis supports master-slave replication

one server is master by default

data replicated to any number of slave servers

multiple ways to set server as slave of master Redis node

in redis.conf : slaveof 127.0.0.1 6379

command line : --slaveof 127.0.0.1 6379

Data dump : can generate / insert data using Redis Ruby gem

can use hiredis gem to speed up data insertion

C driver considerably faster than native Ruby driver

so w/hiredis ...

should be used in production environment

huge performance improvement when using pipelined operation

pipelining reduces number of Redis connections required

but building pipelined dataset has overhead of its own

"find balance between resource consumption and performance improvement"

cluster : interface for building simple ad-hoc distributed Redis cluster

unlike master-slave setup, all servers take master (default) configuration

bridging between servers require minor changes to data dump client

need to require redis/distributed file from redis gem

Redis client replaced with Redis :: Distributed

now take list of server URIs instead of one

decrease in performance because more work done by client

can validate key stored on separate servers

by attempting to retrieve same key from each server

bloom filter : simply never perform query for data that won't return any data

good way to improve performance of any data retrieval system

short-circuiting query for data known to have never been inserted

can return false-positive, cannot ever return false-negative

like bloom filter with HBase

Database - P106

939 - send to 57

Redis

bloom filter: succeed at discovering nonexistence by converting value to very sparse sequence of bits

and comparing sequence to union of every value's bits

each time new value added

OR against current bloom filter bit sequence

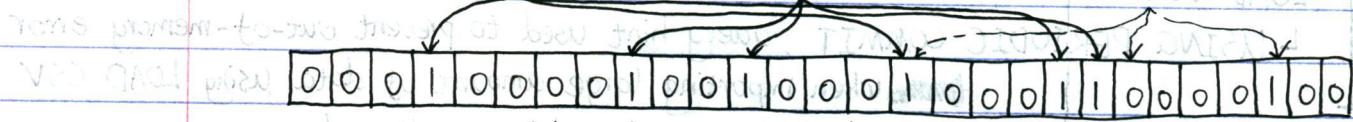
when check whether value already in the system

these operations perform AND against bloom filter bit sequence

because if any true bit not true in bloom filter corresponding bucket

value never added

a b c (not added)



when using bloomfilter-rb Ruby gem created by Ilya Grigorik

concept transferable to any programming language

BIT: bloomfilter function by flipping certain bit in sparse binary field

SETBIT <key> <offset> <value>

set/clear bit at offset in the string value

bet set/cleared depending on <value>, either 1 or 0

if <key> not exist, new string created

string grown to make sure can hold bit at <offset>

<offset> required to be no less than 0, and small than 2^{32}

GETBIT <key> <offset> return bit value at offset in the string value

if <key> not exist, assumed to be empty string

if <offset> beyond string length

string assumed to be contiguous space with 0s

commonly used when faster to flip bit to describe option

BITCOUNT <key> [<start> <end>]

count number of set bit (population counting) in string

by default all bytes contained in string examined

possible to specify counting operation only in interval

passing additional argument <start> and <end>

if <key> not exist, treated as empty string, return 0

Database - P107

bitfield

BITFIELD

Redis

BIT : BITFIELD <key> [GET <type> <offset>] [INCRBY <increment>]

[SET <type> <offset> <value>]

[INCRBY <type> <offset> <increment>]

[OVERFLOW WRAP | SAT | FAIL]

treat Redis string as array of bit

capable of addressing specific integer fields
of varying bit width and arbitrary (non necessary) aligned offset

GET <type> <offset> : return specified bit field

SET <type> <offset> <value> : set specified bit field and return old value

INCRBY <type> <offset> <increment> : increment specified bit field and return new value

OVERFLOW WRAP | SAT | FAIL

sub command only changes behavior of successive INCRBY

WRAP : wrap around, both with signed / unsigned integer

unsigned : modulo maximum value integer can contain

Signed : overflow restart towards most negative value

underflow restart towards most positive value

used if not otherwise specified

SAT : use saturation arithmetic

underflow set to minimum integer value

overflow set to maximum integer value

FAIL : no operation performed on overflow / underflow detected

corresponding return value set to NULL, nil

<type> : prefixing with "i" for signed integer, "u" for unsigned integer

with number of bits of integer type

support up to 64-bit for signed integer

up to 63-bit for unsigned integer

<offset> : if number without any prefix specified

used as 0-based bit offset inside string

if number prefixed with character "#"

specified offset multiplied by integer type width

Database - P108

Redis ~~操作~~ BIT:BITOP <operation> <dest key> [<key>] ~~操作~~

perform bitwise operation between multiple keys

Store result in the destination key

operation : support AND / OR / XOR / NOT

特别地对于 NOT, 只传入一个 key

即有 BITOP NOT <dest key> <key>

if operation performed between strings having different lengths

strings shorter than longest treated as zero-padded up to length of longest

return size of string stored in destination key

BITPOS <key> <bit> [<start>] [<end>]

find the first bit set/clear in string

by default all bytes contained are examined

specified interval passed by additional argument <start> and <end>
interpreted as range of byte rather than range of bit

if find set bit in empty string / string with all 0s, return -1

if find clear bit in string with all 1s

return first bit not part of string on the right

strength : obvious strength of Redis is speed

Redis allow store complex value (list / hash / set), unlike other key-value store

retrieve value based on operation specific to datatype

Redis' durability option allow trade speed for data safety

built-in master-slave replication is great way to ensure durability

without requiring slowness of syncing append-only file on disk on every

replication great for high-read system such as cache

weakness : fast largely because reside in memory

inherent durability problem for main memory database

shut down before snapshot occur will lose data

risk with playing back expiry value even enable append-only file on disk

not support dataset larger than available RAM, partial limitation on size

although Redis clustering available and can help with RAM requirement

must roll own client to do so