

Operating System - P68

dynamic priority adjustment (动态优先级调整), to improve performance
don't always know in advance which jobs are CPU bound or I/O bound
behavior may change at different phases of execution
usually want to give higher priority to I/O bound and interactive job
dynamically adjust the priority based on current process behavior
process use entire time quantum without blocking
can be adjusted downward slightly in priority
repetition will gradually decrease priority
process block before using entire time quantum
can be adjusted upward based on fraction of time quantum used
repetition will gradually increase priority

SPN (shortest process next scheduling, 最短进程优先)

predict the execution time of an interactive command
get optimal response time

prediction value is used to run the estimated shortest process next

based on previous times using a decaying weighted average
 α ($0 < \alpha < 1$) : weighting factor

T_i : measured execution time for execution i] $i \in N$

w_i : weighted average time after execution i

$$w_0 = T_0 \text{ (initial value)}$$

$$w_i = \alpha w_{i-1} + (1-\alpha)T_i, i > 0$$

smaller α : recent behavior have more influence

larger α : past behavior have more influence

combine shortest process next (SPN) with priority scheduling

predict to have shorter execution time with higher priority

guaranteed scheduling (保证调度)

guarantee that each process will receive particular fraction of available CPU time

system tracks CPU usage since process creation

compute ratio of time actually used to time guaranteed

run process with lowest ratio

recompute periodically (each scheduler run)

Operating

System - P69

lottery scheduling (彩票调度), scheduler pick next process to run randomly
each process is assigned probability of being chosen
probability assignment similar to priority system

fair-share scheduling (公平分享调度), similar to guaranteed scheduling
except that CPU allocation is per user instead of per process

time plays essential role in scheduling in real-time system

category : hard real time (硬实时)

soft real time (软实时)

periodic and aperiodic

schedulable (可调度的) : m is the number of periodic events

C_i is CPU time required for event i in one day

P_i is the length of the period of event i

schedulable $\Leftrightarrow \sum_{i=1}^m C_i/P_i \leq 1$

scheduling policy (调度策略): importance level of each process

how to adjust priority

thread

how to apportion CPU among user, process, kernel

scheduling mechanism (调度机制): implementation of the scheduling policy

memory manager (内存管理), keep track of occupied/available main memory

cache management done in hardware

register management done by compiler

allocate memory to process when need

tell process which part of free memory to use

deallocate memory when process done using

mark the memory area as available

relocation problem (重定位问题)

program assume that the first instruction at memory address 0

if program placed at a location other than was assumed

all address based calculation fail

relative address used without any base register dependency

Operating

System - P70

static relocation (静态重定位)

before starting execution, modify all address in program by adding the actual start address of the program
complicated since not all numbers are address

process address space (进程地址空间), set of addresses process can use

static relocation complicated in software

disconnect link between address in program instruction and physical memory
each process will have own address space

address x in one process different from address x in other

logical (virtual) address: address generated by program

each process created assuming: the only one running

program or object placed at physical memory address 0

dynamic relocation (动态重定位)

memory management unit (MMU), hardware component doing address translation

simple MMU has two registers: Base, Limit

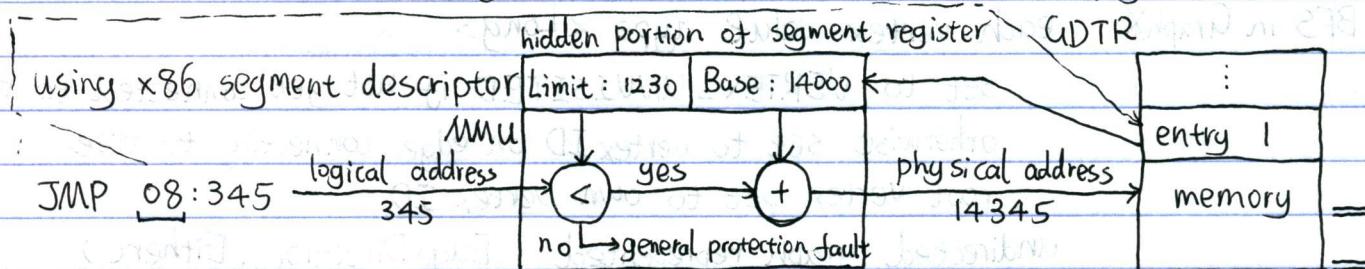
when program runs: load Base with start address of program in physical memory

load Limit with size of program

when refer to memory address m: base + offset

MMU check if m less than the value in Limit

add Base to m to determine the actual physical address



memory allocation (内存分配), main memory usually divided into two partition

resident operating system, usually held in low memory with interrupt vector
user process held in high memory

keep track of free memory: quick allocation / de-allocation

allocate memory area: contiguous / non-contiguous

Operating

System - P71

tracking memory usage with bitmap (位图)

divide memory into allocation units or frames

memory always allocated in multiple of allocation units

tracking memory usage with list (链表)

free, start, size



compaction (压缩), all contiguous free areas less than frames process need

Compact memory by moving jobs
potentially time-consuming

address binding, giving each instruction/variable one memory address

program may be loaded in any part of physical memory

compiler : bind symbolic address to relocatable address

loader : bind relocatable address into absolute address

address binding done in different stage

compile time: absolute code can be generated if memory location known
must recompile code if starting location change

load time: address assigned while loading program into memory

must generate relocatable code at compile time

execution time: binding delayed until run time

if process can be moved during execution

need hardware support for address map

multistep processing

compile
time

Source program

compiler / assembler

object module

other object module

load
time

load module

linkage editor

system library

execution
time

dynamically loaded
System library

dynamic
linking

in-memory binary memory image

Operating

System - P72

EA = physical address

logical address (逻辑地址), generated by CPU, referred to virtual address for the CPU, program is contiguous array of memory address

physical address (物理地址), seen by the memory unit, real address on hardware

logical address : same in compile-time binding + static

and load-time address - binding scheme

can be different in execution-time scheme

new symbols in C++ longjmp stack

linear address (线性地址), one level in x86 protected mode

generated by the CPU using the segment table entry

referred to virtual address instead of logical address

equivalent to physical address when not using Paging

MMU (memory management unit, 内存管理单元), hardware :: logical address → physical address
user program deal with logical address
never see real physical address

relocation register used to protect user processes from each other

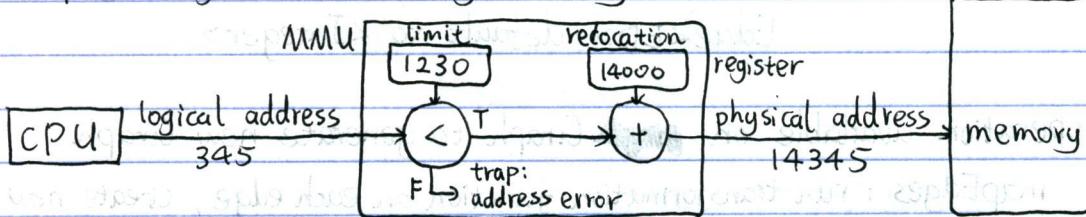
and from changing operating-system code and data

register contents: base (relocation) register, value of smallest physical address

limit register, range of logical address

logical address: must less than limit register

map the logical address dynamically



partition: memory divided into partition, at most one program in one partition

fixed partition: partition size fixed

variable partition: OS monitors parts of memory

which are occupied and which are available

does not do direct bus grafting no bus

does no bus grafting no bus

Operating System - P73

allocation hole (free blocks of memory) of various size
Scattered throughout memory
process allocated memory from contiguous block of memory
large enough to accommodate
OS maintain information about allocated partition and free partition (hole)

dynamic storage allocation

first fit (首次适配): first hole found

best fit (最佳适配): the smallest hole of size large enough
result in smallest leftover hole

worst fit (最差适配): the largest hole

result in largest leftover hole

external fragmentation (外部碎片)

total memory space exist to satisfy a request, but not contiguous
free memory space get broken into little fragmentation

as process loaded and removed

internal fragmentation (内部碎片)

allocated memory may be slightly larger than requested memory
when allocate fixed-size block

extra memory is internal to block and never used

compaction (紧缩) reduce external fragmentation

shuffle memory content to place all free memory together in one large block
possible only if relocation dynamic, done at execution time
use non-contiguous allocation

Paging / Segmentation

part of process can be different parts of memory

32-bit Logical (virtual) address space

text executable instruction	data global variable	heap dynamically allocated memory	stack	kernal address space cannot touch
0x00000000	process address space		0xBFFF	0xFFFFFFFF

Operating System - P74

paging

frame : divide physical memory into fixed-size allocation unit

page : divide logical memory into unit of same size as frame

keep track of all free frames

for process of n frames, find n free frames and load process in the frame

page table : set up to translate logical to physical address

32-bit (14 byte) page table entry

12:31	frame number (20-bit)
9:11	000
8	global page
7	0
6	page modified (dirty bit)
5	page accessed (read bit)
4	disable caching
3	write-through caching
2	accessible by all
1	write allowed
0	page present (present bit)

address translation scheme for given logical address space of 2^m and size of

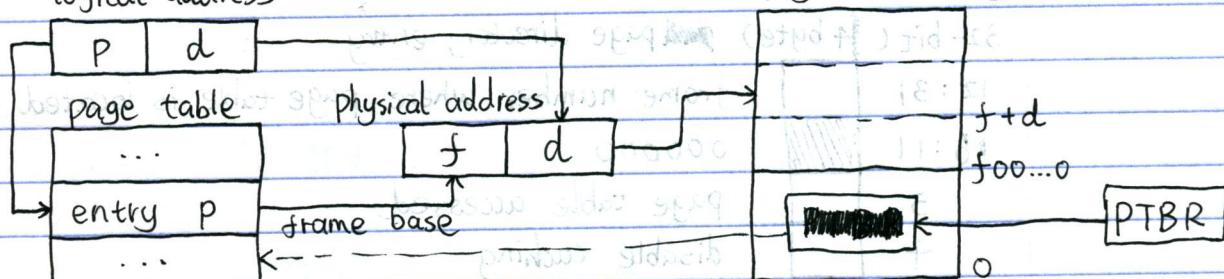
logical address divided into $P(m-n\text{-bit})$ and $d(n\text{-bit})$

page number (P) : use as index into page table

page offset (d) : get base address of frame corresponding to page

page offset (d) : add to base address to determine physical address

logical address



PTBR (page table base register) : loaded with start physical address of page table

page table kept in main memory, PTBR point to page table

Operating

System - P75

TLB

(translation look-aside buffer), 转换检测缓冲区)

special fast-lookup hardware cache

associative, high speed memory

make faster: every data/instruction access require two memory addresses

one for page table, one for data/instruction

high speed parallel search

input: page number

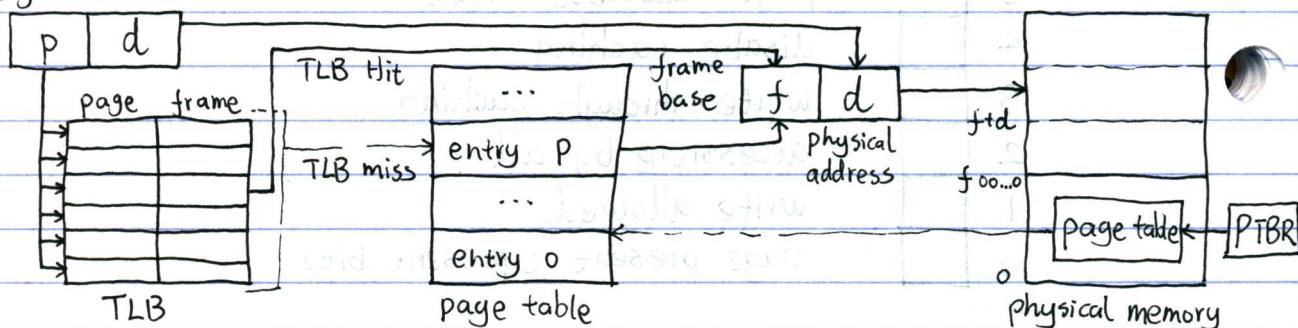
lookup: parallel lookup of associative array for page number

output: frame number / TLB miss

frame reference must be obtained from page table
and the TLB must be updated

TLB flush/update necessary during process switch

logical address



two-level paging, logical address broken down to 3 parts

- | | |
|----------------|---|
| P ₁ | the page directory entry |
| P ₂ | the index into the page table pointed by P ₁ |
| d | add to base address to determine the physical address |

32-bit (4 byte) page directory entry

12:31 frame number where page table is located

11:11 00000

5 page table accessed

4 disable caching

3 write-through caching

2 accessible by all

1 write allowed

0 page table present

Operating

System - P76

x86 paging model	32-bit paging	directory 10-bit directory 10-bit	table 10-bit offset + 22-bit	offset + 12-bit	4KB page 4MB page
------------------	------------------	--------------------------------------	---------------------------------	-----------------	----------------------

PAE paging	directory pointer	2-bit 2-bit	directory 9-bit directory 9-bit	table 9-bit offset 22-bit	offset 12-bit	4KB page 4MB page
---------------	----------------------	----------------	------------------------------------	------------------------------	---------------	----------------------

X86-64 Paging	PML4 PML4 PML4	Dir. Ptr. Dir. Ptr. Dir. Ptr.	directory directory directory	table offset 21-bit offset 30-bit	offset 12-bit	4KB page 2MB page 1GB page
---------------	----------------------	-------------------------------------	-------------------------------------	---	---------------	----------------------------------

inverted page table (倒排页表), one entry for each frame in memory

each entry consist of : page stored in the frame

process that own the page

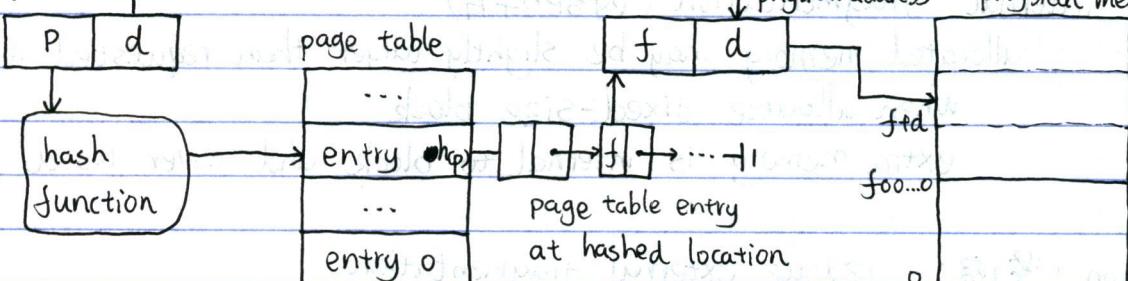
decrease memory needed to store each page table

only one page table in system

increase time needed to search the table when page reference occurs

use hash table to limit the search to few page table entry

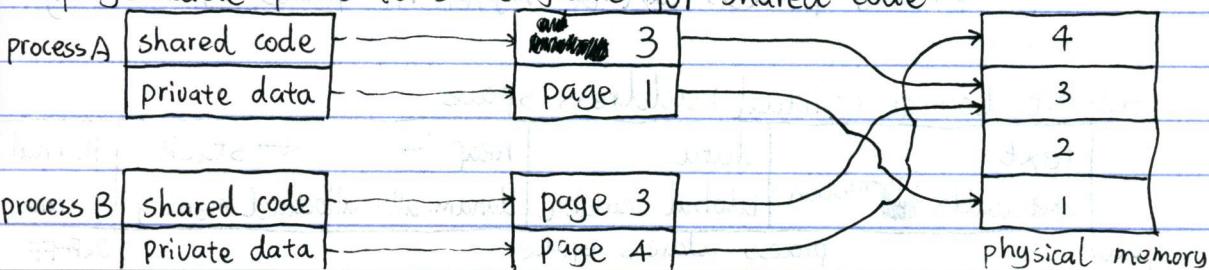
logical address → physical address → physical memory



shared page, shared code : one copy of read-only (reentrant) code
shared among processes

private data : each process keep separate copy of data

page table point to same frame for shared code



Operating

System - P7.7

virtual memory, only part of program need to be in memory for execution

keep only some pages of the process in physical memory

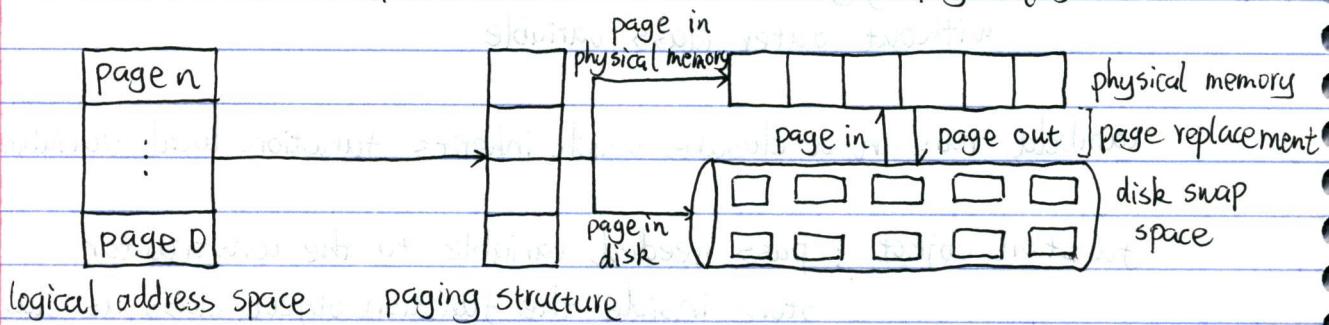
other kept in special disk space (swap space)

Using paging structure to indicate certain page not in physical memory

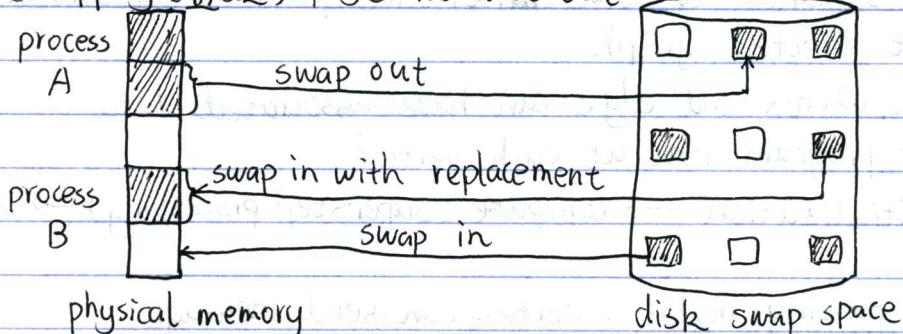
bring page in from swap space when needed

Swap page out to swap space when low in memory

process used up maximum number of allowed physical frame



Swapping (置换) page in and out



page fault present bit : associated with each page table entry.

1: in-memory, 0: not-in-memory

during address translation, present bit in page table

1: access frame 0: page fault

mean trying to access page not in memory

can be because page not mapped / in swap space

Exception thrown if page fault happen

when page fault address unmapped

Operating system must decide allocate page or abort (segmentation fault)

Os can store swap space location in page table entry when page swapped out

Operating

System - P78

page fault

find free frame (with/without replacement)

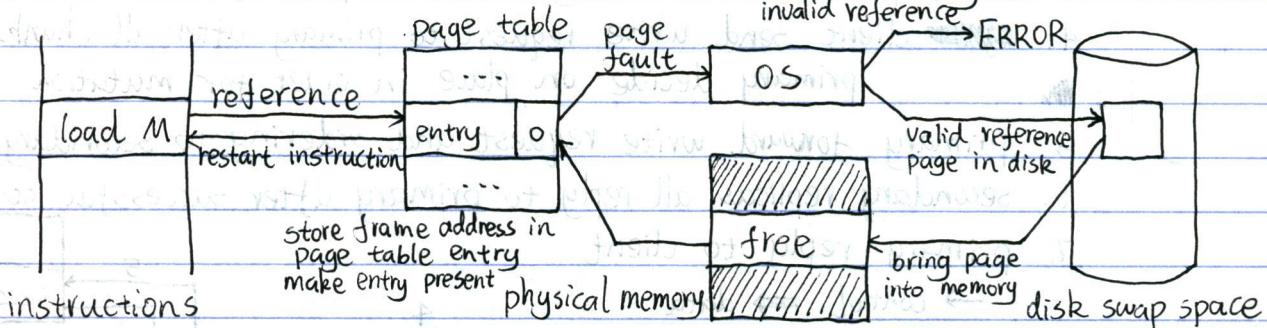
read in desired page into allocated frame

modify page table entry

restart the instruction ~~cause~~ cause the page fault

Program counter point to ~~the~~ faulting instruction

hardware must be able to do this automatically



no free frame, find some frame in memory not really in use, then swap out

need page replacement algorithm

should result in minimum number of page fault

same page may be brought into memory several times

basic page replacement

find location of desired page on disk

find free frame : if there is a free frame, then use

if no free frame, use page replacement algorithm

random, FIFO, LRU, etc. to select victim page

write victim page to disk if page modified

update victim page's page table entry

to indicate page in swap space (can store location here as well)

bring the desired page into the newly free frame

restart user process

reference string, string of page number accessed by process during execution

useful in computing number of page fault will occur

when particular replacement algorithm used