

Operating System - P90

requirement
mutual exclusion: if one process executing in critical section
no other process can be executing in critical section
progress: decision about which enter critical section next
cannot be postponed indefinitely
current process in remainder section cannot participate
bounded wait: limit on how many processes
will first gain access to critical section before given process

synchronization hardware: uniprocessor system can disable interrupt
currently running code will execute without preemption
require elevated privilege level
generally too inefficient on multiprocessor system
OS using this not broadly scalable
modern machine provide special atomic hardware instruction
either test memory word and set value
or swap content of two memory words
atomic: operation guaranteed to happen without interruption
implementation on x86 of hardware locking
XCHG: perform atomic swap (exchange) of two operands
implicitly lock memory location for all processes

semaphore (信号量), integer variable can be initialized to non-negative number
can only be accessed through two standard atomic operations
down/up (wait/signal, wait/release, acquire/release)

busy waiting (忙等待), process waiting to enter critical section in while loop of down operation
spinlock (互斥锁): semaphore implementing for busy waiting
can be removed by associating waiting queue with each semaphore
process block if semaphore value not positive
happen after process added to waiting queue
process restarted when another process execute signal (up)
process to restart chosen from the queue

Operating System - P91

semaphore implemented as system call
OS disable interrupt when running up/down
implemented in thread runtime system
must use one of software synchronization method
to ensure another up/down operation not initiated

binary semaphore (二元信号量) : semaphore value 0/1

counting semaphore (计数信号量) : any other semaphore

starvation (饥饿), process keep waiting to enter critical section

can happen if process waken up from waiting queue

from waiting queue in last-in-first-out order

priority inversion (优先级反转), lower priority process has lock on resource required by higher priority process

mutex lock (互斥锁), implement in kernel or in thread runtime system

ownership: only process own the lock can unlock

modify mutex_lock to remember the process own the lock
run the function without blocking

modify mutex_unlock to execute body only called by owner

futex (fast hybrid mutex) (快速混合互斥锁)

acquiring lock using TestAndSet/Swap much faster than making system call

built-in CPU operation don't provide queuing, require busy wait

implement mutual exclusion operation
provide best features of both

initially use CPU operation to try to get lock

after certain amount of time or number of unsuccessful tries

make kernel call to be put on wait queue

has advantage of not needing kernel call most of time

not trying up the CPU when there is one longer wait

Operating System - P92

condition variable : condition used to block process in Semaphore and mutex lock is simple
(条件变量) no check on any variable

wait(cc) : process that invoke is suspended

signal(cc) : resume one of process that invoked wait(cc) (if any)

usage : user code obtain lock on shared variable

user code check condition on variable

release lock on shared variable

call wait/signal depending on result of condition check

monitor (管程), high-level abstraction provided by language (Java)

provide convenient and effective mechanism for process synchronization

only one process may be active within monitor at one time

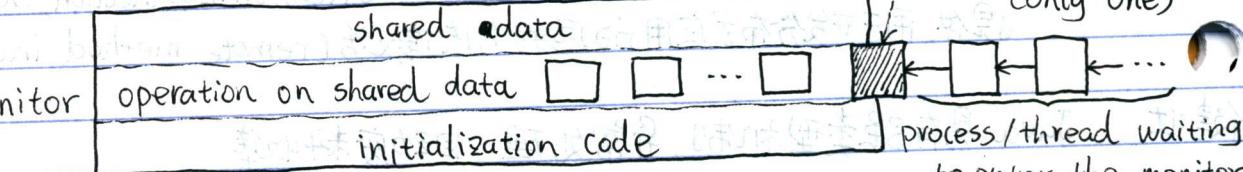
process/thread currently executing inside monitor (only one)

shared data

operation on shared data

initialization code

process/thread waiting to enter the monitor



condition variable : construct used to suspend execution while waiting for condition to become true

difference from semaphore

wait() always block the calling process

wait on semaphore may not, semaphore value maybe positive

signal() has no effect if no process waiting on the condition variable

signal on semaphore always increment value

and wake up any waiting process

process ready to resume because of signal()

must first acquire lock on monitor

Java monitor, Java class may have synchronized member function

only one thread may execute any synchronized function at one time

other thread block

use wait() / notify() instead of condition variable

special version for synchronized function, not subject to race condition

Operating System - Pg 3

reader/writer problem (读者/写者问题)

multiple processes may have simultaneous read-only access to data
only single process at one time may have write access
write access also blocks all readers

need solution that prevent deadlock and starvation

shared resource (memory area/file) protected by critical section
any number of readers may simultaneously be in critical section
when writer need access, must have exclusive access

no reader/other writer have access

writer may be subject to starvation
if there is constant stream of readers

dead lock (死锁), a set of blocked processes

each holding resource and waiting to acquire resource held by another

system model: resource type R_1, R_2, \dots, R_m

CPU cycle, I/O device, mutex lock, semaphore
resource cannot be preempted

each resource type R_i has E_i instances

each process utilizes resource

request: wait if request to acquire not granted immediately

use: can operate on the resource

release: no longer operate on the resource

: deadlock can arise only when all conditions hold simultaneously

mutual exclusion

only one process can use one instance at one time

hold and wait

process holding at least one resource can request additional resource

which may be held by other process

not preemption

resource can be released only voluntarily

circular wait

exist one set $\{P_0, P_1, \dots, P_n\}$ of waiting processes

P_i waiting for resource held by $P_{(i+1)\%n+1}$

Operating System - Pg 4

resource-allocation graph, vertices partitioned into two types

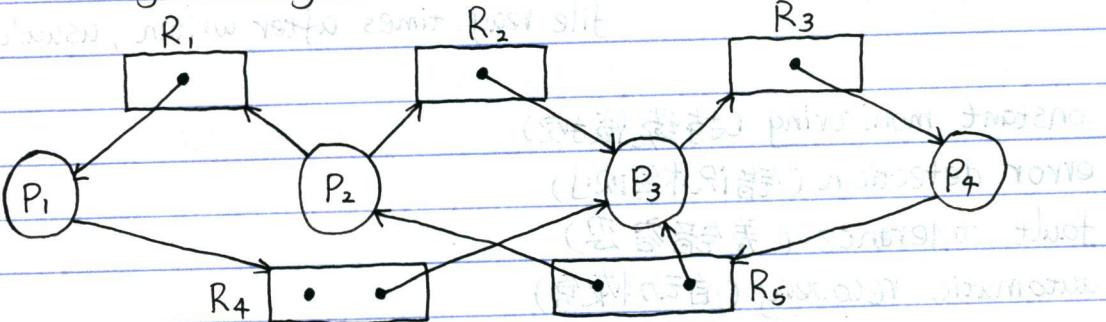
consisting of all processes in the system

consisting of all resource types in the system

edge partitioned into two types

request edge, directed edge (P_i, R_j) , P_i wait R_j

assignment edge, directed edge (R_j, P_i) , P_i hold R_j



cycle not necessarily a deadlock

graph contain no cycle \rightarrow no deadlock

graph contain cycle

only one instance per resource type \rightarrow deadlock

if several instances per resource type, then possibility of deadlock

single instance resources in the system

deadlock exists \Leftrightarrow cycle in the resource-allocation graph

handling: ignore and pretend that deadlock never occur in the system

used by most operating system (Linux)

allow system to enter deadlock state and recover

deadlock detection

ensure that system will never enter deadlock state

deadlock avoidance: temporarily withhold resource if lead to deadlock

deadlock prevention: design system so deadlock is not possible

detection: allow system to enter deadlock state

examine the state of system to determine if deadlock occurred

new request granted / periodically / CPU utilization drop below threshold

recovery scheme: policy on what to do if deadlock occurred

Operating

System - P95

possible recovery scheme

abort all deadlocked processes

abort one process at one time until deadlock cycle eliminated

abort process that not deadlocked

but hold needed resource to break deadlock

preemption : take resource away from another process

rollback : return to some safe state (when resource not needed)

and resume process from that state

much easier to terminate and restart

avoidance

require each process declare maximum of each resource type needed

deadlock - avoidance algorithm

dynamically examine resource - allocation state

ensure there can never be circular-wait condition

resource - allocation state : number of available / allocated resource

maximum demand of process

safe state \exists safe sequence $\langle P_1, P_2, \dots, P_n \rangle$ of all processes in the system

$\forall P_i$, resource P_i may still request can be satisfied by

currently available resource

and resource held by all P_j with $j < i$

unsafe : no safe sequence exist in the system

system must decide if immediate allocation leave system in safe state

when process request available resource

execution in safe state

P_i can wait until all $P_j (j < i)$ finish if resource needed not immediately met

then P_i can obtain needed resource, execute, return allocated resource, terminate

P_1, P_2, \dots, P_i terminate, P_{i+1} may obtain resource needed

system in safe state \rightarrow no deadlock

system in unsafe state \rightarrow possibility of deadlock

avoidance : ensure system never enter unsafe state

safe
unsafe
deadlock

Operating

System - Pg 6

avoidance algorithm : single instance of one resource type

resource allocation graph

multiple instances of one resource type

Banker's algorithm

used when resource can have multiple instances

each process must declare

maximum instance of each resource type needed

when process request resource may have to wait

when process get all resources needed

must return in finite amount of time

need nxm matrix, $\text{Need}[i,j] = k$

process P_i need k more instances of R_j to execute

similar in concept to deadlock detection, except: process

look ahead to maximum of each resource needed by each

spare down the allocation after run algorithm before assigning resource

check if process still in safe state

disadvantages: Banker's algorithm and other algorithm not generally practical

require knowledge of entire future work load of system

and all resources that will be required by process

deadlock prevention, can be incorporated in system design

design to prevent one or more deadlock conditions

design decision place limitation on behavior of process

typically implemented internally for the OS

difficult to enforce limitation on user processes

eliminating mutual exclusion

not allow process to reserve resource

virtualize access to all resources

eliminating hold and wait

require process to request all resources needed as a package

not assign any resource until all resources available

Operating

System - Pg 7

deadlock prevention, eliminating no preemption

possible to implement system policy

allow for preemption of assigned resource instance

System can have mechanism to notify process

not all resource easily preemptible without problem

eliminating circular wait

assign unique number to each resource type

require process to request resource in ascending order

used on CDC 6600 Scope OS. for internal system lock

security problem, computer consist of collection of object, hardware or software

each object has one unique name

and accessed through well-defined set of operation

ensure each object accessed correctly

and only by process allowed to access

secure if resource used /accessed as intended under all circumstance

term intruder (cracker) : attempt to breach security

threat : potential security violation

arising from existence of vulnerability in the system

attack : attempt to breach security

typically by exploiting vulnerability

preserve confidentiality : release of data to unauthorized user never occur

integrity : unauthorized user should not be able to modify any data

availability : system should be serving the purpose it was designed for

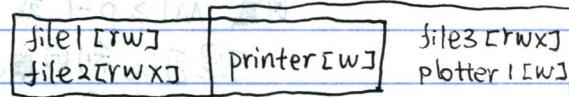
protection domain : set of access-right assigned based on principle of least privilege

access-right = <object-name, right-set>

subset of all valid operation performed on the object

every process run in protection domain

domain 1



Operating

System - Pg 8

domain implementation (UNIX), domain = user-id

domain switch accomplished via file system

each file associated with domain bit (setuid bit)

when file executed and setuid = on

user-id set to owner of file being executed

when execution complete, user-id reset

domain switch accomplished via password

when other domain password provided

sudo command temporarily switch to another user's domain

domain switch via command

sudo command prefix execute specified command in another domain

if original domain has privilege or password given

protection matrix	domain \ object	file 1	file 2	printer	plotter	D1	D2
	domain 1	RW		W			enter
	domain 2		RWX	W	W	enter	

operation "op" must be in cell (i, j)

if process in domain D_i try to do "op" on object O_j

enter operation: domain can switch to another domain

ACL

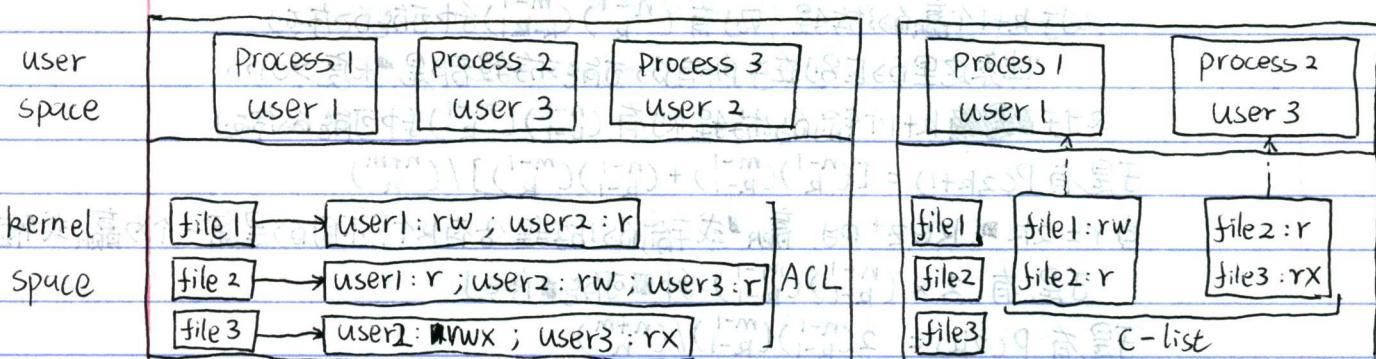
(access control list), each column implemented as access list for one object

resulting per-object list consist of ordered pair <domain, right-set>

defining all domain with non-empty set of access right for object

C-list : (capability list), each process has capability list for domain

list of object together with operation allowed



Operating

System - Pg9

comparison

right check for operation

ACL : check through long list of <domain, right>

C-list : only check capability list of process

revocation : remove right for object

ACL : search access list of object

and remove entry corresponding to the right

can be selective

C-list : search all C-list for object with particular right , then remove

selective removal is difficult

cryptography , constrain potential sender (source) and/or receiver (destination) of message

based on secret key

enable : receipt only by certain destination

confirmation of source

key

string

plaintext (bitstream) → hash algorithm → digest (fixed-size alphanumeric)

attacker can pre-compute digest for possible password and compare

12-bit salt

+ Password → hash algorithm → digest

password

salt* should be random, large number (>12-bit) and different for each stored and accessed for password checking

symmetric encryption

plaintext (bitstream)

encryption algorithm

(garbled bitstream)

secret key

ciphertext

decryption algorithm

attacker know the algorithm and can see ciphertext

obtaining secret key from ciphertext computationally infeasible
public key of receiver

asymmetric encryption

plaintext

encryption algorithm

ciphertext

private key of receiver

decryption algorithm

attacker know algorithm, public key and can see ciphertext

obtaining private key from ciphertext and public key is computationally infeasible

Operating

System - P100

asymmetric encryption (非对称加密) → public-key encryption based on each user having two keys

public key : published key used to encrypt data

private key : key known only to individual user

used to decrypt data

encryption scheme can be made public

without making easy to figure out the decryption scheme

most common : RSA block cipher

efficient algorithm for testing prime number

no efficient algorithm for finding prime factor of large number

digital signature : message + digest (hash) → **encryption algorithm** → private key (sender)

sender
receiver

ciphertext

(sender)

message + d (d' = digest) ← **decryption algorithm** ← public key

message not tampered with during transit if $d = d'$

then sender also authenticated

key distribution delivery of symmetric key is huge challenge

sometimes done out-of-band
asymmetric key can proliferate

public key no secret

asymmetric key distribution need care
man-in-the-middle attack

digital certificate , proof of ownership of public key

trusted party receive proof of identification from user

certificate public key belong to user

user's public key encrypted / signed with trusted party's private key

signature legitimate

certificate authorities are trusted parties

public key included with web browser distribution

vouch for other authorities via digitally signing their keys